

Hieu Nguyen

SID: 26369732

**Group:** I worked with a Tree on this one, but I did all the work.

## 1 Academic Integrity Statement

“I certify that all solutions are entirely in my own words and that I have not looked at another student’s solutions. I have given credit to all external sources I consulted.”

*Hieu Nguyen*

– Hieu Nguyen

## 2 Gaussian Classification

1. Finding the Bayes optimal decision boundary for point(s) at which the posterior probabilities are equal

$$\begin{aligned}
 P(C_1|x) = P(C_2|x) &\iff f(x|C_1) = f(x|C_2) \\
 &\iff \frac{(x - \mu_1)^2}{2\sigma^2} = \frac{(x - \mu_2)^2}{2\sigma^2} \\
 &\iff (x - \mu_1)^2 = (x - \mu_2)^2 \\
 &\iff x - \mu_1 = \mu_2 - x \\
 &\iff x = \frac{\mu_1 + \mu_2}{2}
 \end{aligned}$$

We have the Bayes optimal decision boundary:  $x = \frac{\mu_1 + \mu_2}{2}$

And the corresponding Bayes decision rule is that for any data point  $x \in \mathbb{R}$ , if  $x < \frac{\mu_1 + \mu_2}{2}$ , then  $x$  is classified as class 1; and if  $x > \frac{\mu_1 + \mu_2}{2}$ , then  $x$  is classified as class 2.

2. Supposed the decision boundary for the classifier is  $x = b$ . The Bayes error is given by

$$P_e = P((C_1 \text{ misclassified as } C_2) \cup (C_2 \text{ misclassified as } C_1))$$

which is equivalent to

$$P_e = P((\text{misclassified as } C_1|C_2)P(C_2) + P((\text{misclassified as } C_2|C_1)P(C_1))$$

For the first part,  $P((\text{misclassified as } C_1|C_2)P(C_2)$ ,

$$P((\text{misclassified as } C_1|C_2)P(C_2) = \frac{1}{2} \int_{-\infty}^b \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu_2)^2/(2\sigma)^2} dx$$

Similarly, we have,

$$P((\text{misclassified as } C_2|C_1)P(C_1) = \frac{1}{2} \int_b^{\infty} \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu_1)^2/(2\sigma)^2} dx$$

Therefore, we have

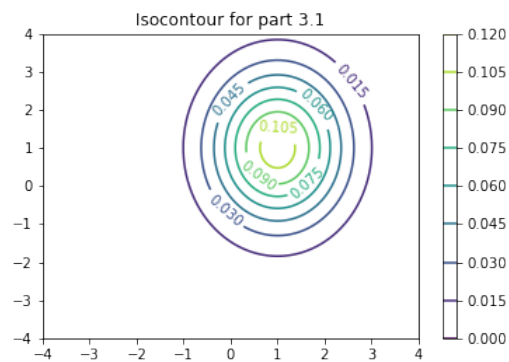
$$\begin{aligned}
 P_e &= P((C_1 \text{ misclassified as } C_2) \cup (C_2 \text{ misclassified as } C_1)) = \\
 &\frac{1}{2\sqrt{2\pi}\sigma} \left( \int_{-\infty}^b e^{-(x-\mu_2)^2/(2\sigma)^2} dx + \int_b^{\infty} e^{-(x-\mu_1)^2/(2\sigma)^2} dx \right)
 \end{aligned}$$

3. We want to calculate the optimal decision boundary  $b^*$  that minimize  $P_e(b)$ , given that  $P_e(b)$  is convex for  $\mu_1 < b < \mu_2$

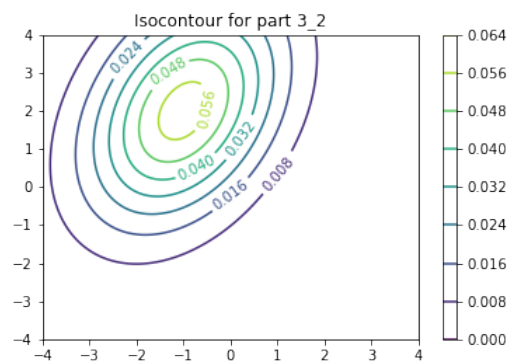
### 3 Isocontours of Normal Distributions

For the code, please see the appendix.

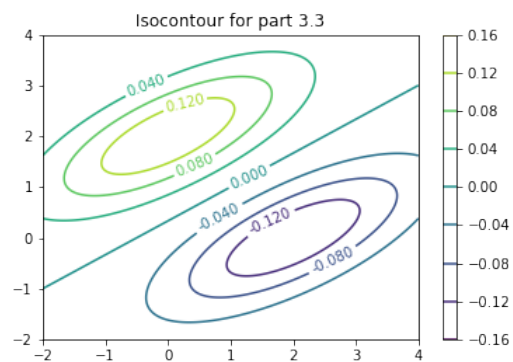
#### 1. Picture




#### 2. Picture



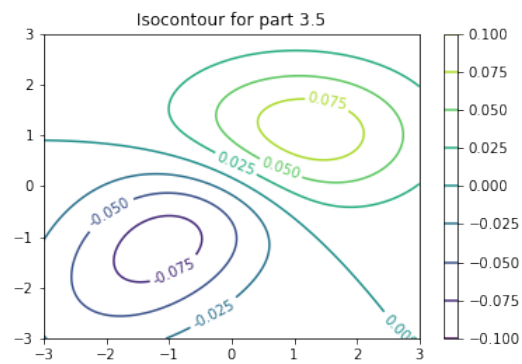
#### 3. Picture



#### 4. Picture



#### 5. Picture



## 4 Eigenvectors of the Gaussian Covariance Matrix

(a) (Please see the appendix for the code) The mean (in  $\mathbb{R}^2$ ) of the sample is  $[3.31998 \ 5.423675]$ .

(b) (Please see the appendix for the code) The  $2 \times 2$  covariance matrix of the sample  $\Sigma$  is

$$\begin{bmatrix} 76.0392 & 45.73699 \\ 45.73699 & 40.59195 \end{bmatrix}$$

(c) (Please see the appendix for the code)


For the eigenvalue of  $\lambda_1 = 107.366$ , the eigenvector is:

$$\begin{bmatrix} 0.8225 \\ 0.5651 \end{bmatrix}$$


For the eigenvalue of  $\lambda_2 = 9.2646$ , the eigenvector is:

$$\begin{bmatrix} -0.5651 \\ 0.8225 \end{bmatrix}$$

(d) Picture



(e) Picture



## 5 Classification and Risk

1. Show that the following policy obtains the minimum risk when  $\lambda_r \leq \lambda_s$

(a) Choose class  $i$  if  $P(Y = i|x) \geq P(Y = j|x)$  for all  $j$  and  $P(Y = i|x) \geq 1 - \lambda_r/\lambda_s$ .

Let  $r : \mathbb{R}^d \rightarrow \{1, \dots, c+1\}$  be a decision rule. And the loss function is defined as follow:

$$L(r(x) = i, y = j) = \begin{cases} 0 & \text{if } i = j \text{ } i, j \in \{1, \dots, c\}, \\ \lambda_r & \text{if } i = c+1, \\ \lambda_s & \text{otherwise,} \end{cases}$$

Then we want to show that, for an arbitrary decision function,  $f$ , regardless of the combination of true label,  $r(x)$ , and  $f(x)$ , we have  $R(r(x)|x) \leq R(f(x)|x)$ .

$$R(r(x) = i|x) = \sum_{j=1}^c L(r(x) = i, y = j)P(Y = j|x) = \lambda_s \sum_{j \neq i} P(Y = j|x) = \lambda_s(1 - P(Y = i|x))$$

- For  $f(x) \neq c+1$ , we have that  $R(f(x) \neq c+1|x) = \lambda_s(1 - P(Y \neq c+1|x))$ , then we can conclude that

$$\rightarrow R(r(x) = i|x) \leq R(f(x) \neq c+1|x)$$

- For  $f(x) = c+1$ , we have that  $R(f(x) = c+1|x) = \lambda_r$ , then we can also conclude that

$$\rightarrow R(r(x) = i|x) \leq R(f(x) = c+1|x)$$

because  $P(Y = i|x) \geq 1 - \lambda_r/\lambda_s \rightarrow \lambda_r \geq \lambda_s(1 - P(Y = i|x))$  which is resulted in the expression above. because  $P(Y = i|x) \geq P(Y \neq c+1|x)$



(b) Choose doubt otherwise.

If  $r(x) = c + 1$ , similar to reasoning above, we have  $R(r(x) = c + 1|x) = \lambda_r$ . Consider an arbitrary  $f(x) \neq c + 1$ , we have

$$R(f(x) \neq c + 1|x) = \lambda_s(1 - P(Y \neq c + 1|x))$$

And since we choose doubt otherwise, we fail to satisfy the conditions of the previous case, then  $P(Y \neq c + 1|x) \leq 1 - \lambda_r/\lambda_s$ . This lead to the conclusion that

$$R(r(x) = c + 1|x) < R(f(x) \neq c + 1|x)$$

- 2.
- What happens if  $\lambda_r = 0$ ? We have  $P(Y = i|x) = 1$  or  $P(r(x) = i|x) = 1$ . So we can choose class  $i$  to classify  $x$  if we know for sure (probability of 1), otherwise, we can choose doubt.
  - What happens if  $\lambda_r > \lambda_s$ ? Then we have  $1 - \lambda_r/\lambda_s < 0$ , this means we should choose class  $i$  (labeled  $1, \dots, c$ ) to classify  $x$  to attain the highest chance for optimal classification.
  - Intuitively, when  $\lambda_r = 0$ , it is better to choose doubt since there are no risks for doing so (with  $\lambda_r = 0$ ); and when  $\lambda_r > \lambda_s$ , it is consistent with the intuition that it's better to classify  $x$  by choosing the class  $i$  that provides the best chance for a correct classification, rather than choosing doubt with higher risk.

## 6 Maximum Likelihood Estimation and Bias

(a) The likelihood function is

$$\mathbf{L}(\mu, \sigma; x_1, x_2, \dots, x_n) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi} \sqrt{\sigma^2/i}} e^{-\frac{1}{2} \frac{(x_i - \mu)^2}{\sigma^2/i}}$$

Then the log-likelihood function is

$$l(\mu, \sigma, \mathbf{X}) = -\frac{n}{2} \ln(2\pi) + \frac{1}{2} \sum_{i=1}^n \ln i - n \sum_{i=1}^n \ln \sigma - \frac{1}{2\sigma^2} \sum_{i=1}^n i(x_i - \mu)^2$$

Looking at the partial derivative of the parameters to find the estimator,

$$\begin{cases} \frac{\partial l}{\partial \mu} = \frac{1}{\sigma^2} \sum_{i=1}^n i(x_i - \mu) = 0 \rightarrow \sum_{i=1}^n i(x_i - \mu) = 0 \\ \frac{\partial l}{\partial \sigma} = -\frac{n}{\sigma} + \frac{1}{\sigma^3} \sum_{i=1}^n i(x_i - \mu)^2 = 0 \rightarrow \sigma^2 = \frac{1}{n} \sum_{i=1}^n i(x_i - \mu)^2 \end{cases}$$

$$\begin{cases} \sum_{i=1}^n i x_i = \mu \sum_{i=1}^n i \rightarrow \sum_{i=1}^n i x_i = \mu \frac{n(n+1)}{2} \rightarrow \mu = \frac{2}{n(n+1)} \sum_{i=1}^n i x_i \\ \sigma^2 = \frac{1}{n} \sum_{i=1}^n i(x_i - \mu)^2 \end{cases}$$

Then we have the MLE estimators for  $\mu$  and  $\sigma$ ,

$$\begin{cases} \hat{\mu} = \frac{2}{n(n+1)} \sum_{i=1}^n i X_i \\ \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n i(X_i - \mu)^2 \end{cases}$$

- (b) • Statement 1: "The MLE sample estimator  $\hat{\mu}$  is unbiased.

$$\begin{aligned}\mathbf{E}[\hat{\mu}] - \mu &= \mathbf{E} \left[ \frac{2}{n(n+1)} \sum_{i=1}^n i X_i \right] - \mu = \frac{2}{n(n+1)} \sum_{i=1}^n i \mathbf{E}[X_i] - \mu \\ &= \frac{2\mu}{n(n+1)} \sum_{i=1}^n i - \mu = \mu - \mu = 0\end{aligned}$$

Therefore, the statement is true and that the MLE sample estimator  $\hat{\mu}$  is unbiased. Note that we use the linearity of expectation above to bring the expectation inside the summation.

- Statement 2: "The MLE sample estimator  $\hat{\sigma}^2$  is unbiased.

$$\mathbf{E}[\hat{\sigma}^2] - \sigma^2 = \mathbf{E} \left[ \frac{1}{n} \sum_{i=1}^n i (X_i - \mu)^2 \right] - \sigma^2 = \frac{1}{n} \sum_{i=1}^n i \mathbf{E}[(X_i - \mu)^2] - \sigma^2$$

Using the fact that  $\mathbf{E}[X^2] = \mathbf{Var}[X] + (\mathbf{E}[X])^2$ , we have

$$\begin{aligned}\frac{1}{n} \sum_{i=1}^n i \mathbf{E}[(X_i - \mu)^2] - \sigma^2 &= \frac{1}{n} \sum_{i=1}^n i [\mathbf{Var}[(X_i - \mu)] + (\mathbf{E}[X_i - \mu])^2] - \sigma^2 \\ &= \frac{1}{n} \sum_{i=1}^n i [\mathbf{Var}[X_i] + 0] - \sigma^2 = \frac{1}{n} \sum_{i=1}^n i \sigma_i^2 - \sigma^2 = \left( \frac{1}{n} \sum_{i=1}^n i \cdot \sigma^2 / i \right) - \sigma^2 = 0\end{aligned}$$

Therefore, the statement is true and that the MLE sample estimator  $\hat{\sigma}^2$  is unbiased.

## 7 Covariance Matrices and Decompositions


- (a)  $\hat{\Sigma}$  is not invertible if and only if there exists a hyperplane in  $d - 1$  dimensions such that all the points can lie on that plane. In the context of linear algebra, this is similar to  $\hat{\Sigma}$  not having full rank.  $\hat{\Sigma}$  is not invertible if and only if all the sample points lie on a common hyperplane in the space that doesn't span all of  $\mathbb{R}^{d \times d}$
- (b) We can make a new matrix by adding a bit of noise to the diagonal so we can make all the eigenvalues of the new covariance matrix positive in order for it to be invertible (i.e:  $\hat{\Sigma} = \hat{\Sigma} + c\mathbf{I}$ ). To avoid biasing the covariance matrix, we can adjust the constant  $c$  to not be too big or too small (certain order of magnitude) than the minimum non-zero eigenvalue of the original covariance matrix. This will avoid blowing up the eigenvalues.
- (c) When  $\mu = 0$ , we have,

$$f(x) = \left( \frac{1}{(\sqrt{2\pi})^d |\Sigma|} \right) e^{-\frac{x^T \Sigma^{-1} x}{2}}$$

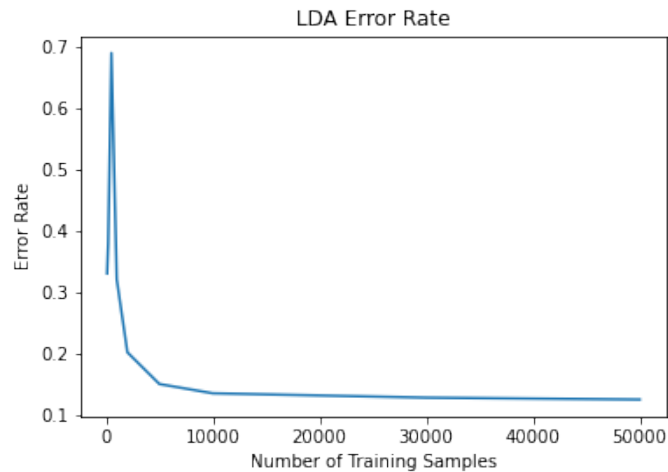
So if we want to maximize  $f(x)$ , we want to maximize  $x^T \Sigma^{-1} x$ . This can be done by finding the largest eigenvalue of  $\Sigma^{-1}$ , and having  $x$  equal to the corresponding eigenvector and normalizing it to have the length one. Similarly, we can minimize  $f(x)$  by letting  $x$  equal to the normalized eigenvector with the smallest eigenvalue of  $\Sigma^{-1}$

## 8 Gaussian Classifiers for Digits and Spam


1. Please see the appendix for the code
2. Below is the visualization for the covariance matrix of digit 0. We can see the the diagonal terms have a higher covariance in comparison to the off-diagonal terms. This indicates that the the covariance between each sample points and itself, as well as each sample point and adjacent points, is higher than the covariance between a sample point and a point farther away.



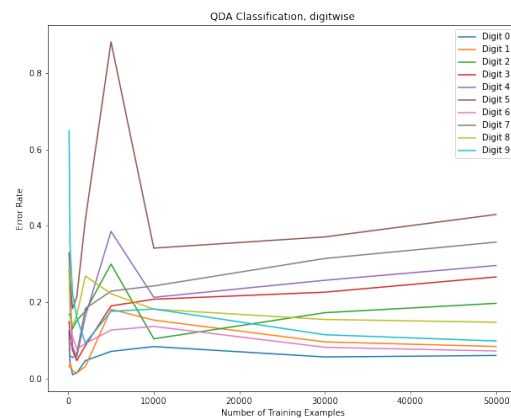
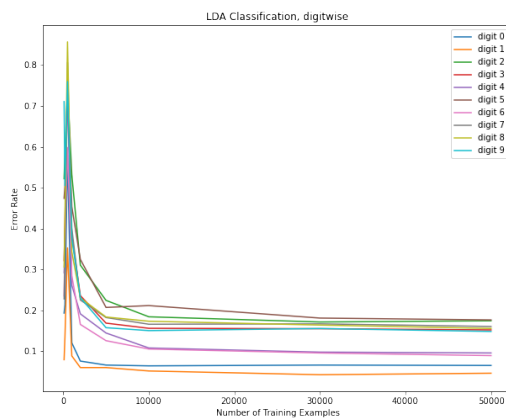
3. (a) Picture



(b) Picture



- (c) LDA performed better than QDA because QDA is prone to overfitting, especially when there are a larger number of free variables
- (d) From the plot, we can see that the it is the best for LDA to classify digit 1 and digit 0 is the best for QDA based on the validation error.



4. **Kaggle Name:** Hieu Tang Nguyen BA

**Kaggle Link:** <https://www.kaggle.com/hieutangnguyenba>

**MNIST Accuracy:** 0.95800

5. **Kaggle Name:** Hieu Tang Nguyen BA

**Kaggle Link:** <https://www.kaggle.com/hieutangnguyenba>

**Spam Accuracy:** 0.80640

## 9 Code Appendix

### 9.1 Constants, packages, etc.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 # ... etc. This is an environment where you can enter code.
5 # You could also just include screenshots using
6 # \includegraphics[options]{image name}
```

### 9.2 Isocontours of Normal Distributions

```
1 import numpy as np
2 from scipy import stats
3 import matplotlib.pyplot as plt
4
5 # 3.1
6
7 x = np.linspace(-4, 4, 1000)
8 y = np.linspace(-4, 4, 1000)
9
10 X,Y = np.meshgrid(x, y)
11
12 pos = np.array([Y, X]).T
13
14 z = scipy.stats.multivariate_normal([1, 1], [[1, 0], [0, 2]])
15
16 Z = z.pdf(pos)
17
18 plt.figure()
19 contour = plt.contour(X, Y, Z)
20 plt.clabel(contour, inline=1, fontsize=10)
21 plt.title('Isocontour for part 3.1')
22 plt.colorbar()
23 plt.savefig("3_1.png")
24 plt.show()
25
26 # 3.2
27
28 x = np.linspace(-4, 4, 1000)
29 y = np.linspace(-4, 4, 1000)
30
31 X,Y = np.meshgrid(x, y)
32
33 pos = np.array([Y, X]).T
34
35 z = scipy.stats.multivariate_normal([-1, 2], [[2, 1], [1, 4]])
36
37 Z = z.pdf(pos)
38
39 plt.figure()
40 contour = plt.contour(X, Y, Z)
41 plt.clabel(contour, inline=1, fontsize=10)
42 plt.title('Isocontour for part 3_2')
43 plt.colorbar()
44 plt.savefig("3_2.png")
45 plt.show()
46
47 # 3.3
48
49 x = np.linspace(-2, 4, 1000)
50 y = np.linspace(-2, 4, 1000)
51
```

```

52 X,Y = np.meshgrid(x, y)
53
54 pos = np.array([Y, X]).T
55
56 z1 = scipy.stats.multivariate_normal([0, 2], [[2, 1], [1, 1]])
57 z2 = scipy.stats.multivariate_normal([2, 0], [[2, 1], [1, 1]])
58
59
60 Z = z1.pdf(pos) - z2.pdf(pos)
61
62
63 plt.figure()
64 contour = plt.contour(X, Y, Z)
65 plt.clabel(contour, inline=1, fontsize=10)
66 plt.title('Isocontour for part 3.3')
67 plt.colorbar()
68 plt.savefig("3_3.png")
69 plt.show()
70
71 # 3.4
72
73 x = np.linspace(-2, 4, 1000)
74 y = np.linspace(-2, 4, 1000)
75
76 X,Y = np.meshgrid(x, y)
77
78 pos = np.array([Y, X]).T
79
80 z1 = scipy.stats.multivariate_normal([0, 2], [[2, 1], [1, 1]])
81 z2 = scipy.stats.multivariate_normal([2, 0], [[2, 1], [1, 4]])
82 Z = z1.pdf(pos) - z2.pdf(pos)
83
84 plt.figure()
85 contour = plt.contour(X, Y, Z)
86 plt.clabel(contour, inline=1, fontsize=10)
87 plt.title('Isocontour for part 3.4')
88 plt.colorbar()
89 plt.savefig("3_4.png")
90 plt.show()
91
92 # 3.5
93
94 x = np.linspace(-3, 3, 1000)
95 y = np.linspace(-3, 3, 1000)
96
97 X,Y = np.meshgrid(x, y)
98 pos = np.array([Y, X]).T
99
100 z1 = scipy.stats.multivariate_normal([1, 1], [[2, 0], [0, 1]])
101 z2 = scipy.stats.multivariate_normal([-1, -1], [[2, 1], [1, 2]])
102
103 Z = z1.pdf(pos) - z2.pdf(pos)
104
105 plt.figure()
106 contour = plt.contour(X, Y, Z)
107 plt.clabel(contour, inline=1, fontsize=10)
108 plt.title('Isocontour for part 3.5')
109 plt.colorbar()
110 plt.savefig("3_5.png")
111 plt.show()

```

## 9.3 Eigenvectors of the Gaussian Covariance Matrix

```

1
2 import matplotlib.pyplot as plt
3 import numpy as np

```



```

4
5 np.random.seed(26)
6
7 size = 100
8 mu_x = 3
9 var_x = 9
10 mu_y = 4
11 var_y = 4
12
13 x_sample = np.random.normal(mu_x, var_x, size)
14 y_sample = np.random.normal(mu_y, var_y, size)
15
16 sample_points = np.array([np.array((x, 0.5 * x + y)) for (x, y) in zip(x_sample, y_sample)])
17
18 # (a): compute the mean of the sample
19
20 sample_mean = np.mean(sample_points, axis=0)
21
22 print('Mean of the Sample:')
23
24 print(sample_mean)
25
26 # (b): compute the 2x2 covariance matrix of the sample
27
28 sample_covariance = np.cov(sample_points.T)
29
30 print('2x2 Covariance Matrix of the Sample')
31
32 print(sample_covariance)
33
34 # (c) compute the eigenvectors and eigenvalues of the covariance matrix
35
36 eigenvalues, eigenvectors = np.linalg.eig(sample_covariance)
37
38 print('Eigenvalues of the Covariance Matrix:')
39
40 print(eigenvalues)
41
42 print("-----")
43
44 print('Eigenvectors of the Covariance Matrix:')
45
46 print(eigenvectors)
47
48 # (d): plot 100 data points and eigenvectors
49
50 plt.figure()
51
52 plt.title("Sample Points and The Eigenvectors")
53 plt.xlabel("X1")
54 plt.ylabel("X2")
55 plt.xlim(-15, 15)
56 plt.ylim(-15, 15)
57
58
59 plt.scatter(sample_points[:, 0], sample_points[:, 1])
60
61 x_vector = [sample_mean[0], sample_mean[0]]
62 y_vector = [sample_mean[1], sample_mean[1]]
63 v_1_vector = [eigenvectors[0][0] * eigenvalues[0], eigenvectors[0][1] * eigenvalues[1]]
64 v_2_vector = [eigenvectors[1][0] * eigenvalues[0], eigenvectors[1][1] * eigenvalues[1]]
65
66 plt.quiver(x_vector, y_vector, v_1_vector, v_2_vector, angles="xy", scale_units="xy", scale=1)
67
68
69 plt.savefig("4d.png")
70 plt.show()
71

```

```

72 # (e): plot rotated points
73
74 rotated_points = np.dot(eigen_vectors.T, (sample_points - sample_mean).T).T
75
76 plt.figure()
77
78 plt.title("100 Rotated Sample Points")
79 plt.xlabel("X1")
80 plt.ylabel("X2")
81 plt.xlim(-15, 15)
82 plt.ylim(-15, 15)
83
84 plt.scatter(rotated_points[:, 0], rotated_points[:, 1])
85
86
87 plt.savefig("4e.png")
88 plt.show()

```

## 9.4 Gaussian Classifiers for Digits and Spam

```

1
2
3 import numpy as np
4 import scipy.cluster
5 import scipy.ndimage
6 import matplotlib
7 import matplotlib.pyplot as plt
8 import os
9 from scipy import io
10 import pandas as pd
11 import pprint
12 from scipy.stats import multivariate_normal
13
14
15 # Usage results_to_csv(clf.predict(X_test))
16 def results_to_csv(y_test):
17     y_test = y_test.astype(int)
18     df = pd.DataFrame({'Category': y_test})
19     df.index += 1 # Ensures that the index starts at 1.
20     df.to_csv('submission.csv', index_label='Id')
21
22 for data_name in ["mnist", "spam"]:
23     data = io.loadmat("data/%s_data.mat" % data_name)
24     print("\nloaded %s data!" % data_name)
25     fields = "test_data", "training_data", "training_labels"
26     for field in fields:
27         print(field, data[field].shape)
28
29 # Load all data for MNIST
30 mnist = io.loadmat("data/mnist_data.mat")
31 mnist_X, mnist_y = mnist['training_data'].astype(float), mnist['training_labels']
32 mnist_test_X = mnist['test_data'].astype(float)
33
34 # Load all data for SPAM
35 spam = io.loadmat("data/spam_data.mat")
36 spam_X, spam_y = spam['training_data'].astype(float), spam['training_labels']
37 spam_test_X = spam['test_data'].astype(float)
38
39
40 ##### 8.1
41
42 # Contrast normalization to the training set.
43 cn_mnist_X = []
44
45 for i in range(len(mnist_X)):
46

```

```

47     val = mnist_X[i]/(np.linalg.norm(mnist_X[i])+1e-15)
48     cn_mnist_X.append(val)
49
50 cn_mnist_X = np.array(cn_mnist_X)
51 cn_mnist_X.shape
52
53 # Contrast normalization to the testing set.
54 cn_mnist_test_X = []
55 for i in range(len(mnist_test_X)):
56     val = mnist_test_X[i]/(np.linalg.norm(mnist_test_X[i])+1e-15)
57     cn_mnist_test_X.append(val)
58
59 cn_mnist_test_X = np.array(cn_mnist_test_X)
60 cn_mnist_test_X.shape
61
62 mnist_fitted = {}
63 for i in np.unique(mnist_y):
64     indices = (mnist_y == i).flatten()
65     data = cn_mnist_X[indices]
66     mean = np.mean(data, axis=0)
67     cov = np.cov(data, rowvar = False)
68     mnist_fitted[i] = (mean, cov)
69
70
71 ### 8.2
72
73 indices = (mnist_y == np.unique(mnist_y)[0]).flatten()
74 data = cn_mnist_X[indices]
75 ncov = np.corrcoef(data, rowvar=False)
76 ncov[np.isnan(ncov)] = 0
77 ncov = np.abs(ncov)
78 plt.imshow(ncov, cmap=matplotlib.cm.Greys_r)
79 plt.colorbar()
80 #plt.savefig("8_2.png")
81 plt.show()
82
83
84 ### 8.3
85
86 def split_train_val_data(data, labels, val_size):
87     num_items = len(data)
88     assert num_items == len(labels)
89     assert val_size >= 0
90     if val_size < 1.0:
91         val_size = int(num_items * val_size)
92     train_size = num_items - val_size
93     idx = np.random.permutation(num_items)
94     data_train = data[idx][:train_size]
95     label_train = labels[idx][:train_size]
96     data_val = data[idx][train_size:]
97     label_val = labels[idx][train_size:]
98     return data_train, data_val, label_train, label_val
99
100 class LDA:
101     def __init__(self):
102         self.mu = None
103         self.sigma = None
104         self.Pi = None
105
106     def fit(self, X, y):
107
108         N, labels = len(y), np.unique(y)
109         d = X.shape[1]
110         C = len(labels)
111         mu = np.zeros((C, d))
112         Pi = np.zeros(C)
113         sigma = np.zeros((d, d))
114         for i, label in enumerate(labels):

```

```

115         X_i = []
116         for index, result in enumerate(y):
117             if result == label:
118                 X_i.append(X[index])
119         X_i = np.array(X_i)
120         #print(len(X_i))
121         Pi[i] = len(X_i) / N
122         mu[i, :] = np.mean(X_i, axis=0)
123         sigma += np.cov(X_i.T)*X_i.shape[0]
124     sigma /= N
125     self.mu = mu
126     self.sigma = sigma
127     self.Pi = Pi
128
129     def predict(self, X):
130
131         d = len(self.sigma)
132         self.sigma += 1e-15*np.array(np.eye(d))
133
134         L1 = self.mu.dot(np.linalg.solve(self.sigma, X.T)).T
135
136         L2 = 1/2*np.diag(self.mu.dot\
137                         (np.linalg.solve(self.sigma, self.mu.T)))
138
139         L = L1 - L2 - np.log(self.Pi)
140
141         pred = np.argmax(L.T, axis=0)
142
143
144         return pred
145
146     def accuracy(self, X, y):
147         N = len(y)
148         correct = 0
149         pred = self.predict(X)
150         for i in range(N):
151             if pred[i] == y[i]:
152                 correct += 1
153         return correct/N, pred
154
155 mnist_train_X, mnist_val_X, mnist_train_y, mnist_val_y = split_train_val_data(cn_mnist_X, mnist_y
156                                     ↪ , 10000)
157
158 training_size = np.array([100, 200, 500, 1000, 2000, 5000, 10000, 30000, 50000])
159 lda_accuracy = np.zeros(len(training_size))
160
161 for i, N in enumerate(training_size):
162     lda = LDA()
163     lda.fit(mnist_train_X[:N,:], mnist_train_y[:N])
164     lda_accuracy[i], _ = lda.accuracy(mnist_val_X, mnist_val_y)
165
166 lda_accuracy
167
168 plt.figure()
169 plt.plot(training_size, (1-lda_accuracy))
170 plt.title('LDA Error Rate')
171 plt.xlabel('Number of Training Samples')
172 plt.ylabel('Error Rate')
173 #plt.savefig("8_3_a.png")
174 plt.show()
175
176 class QDA:
177     def __init__(self, a):
178         self.mu = None
179         self.sigma = None
180         self.Pi = None
181         self.a = a

```

```

182
183 def fit(self, X, y):
184     N, labels = len(y), np.unique(y)
185     #print(labels)
186     d = X.shape[1]
187     C = len(labels)
188     sigma = [np.zeros([d, d]) for i in range(N)]
189     mu = np.zeros([C, d])
190     Pi = np.zeros([C])
191     for i, label in enumerate(labels):
192         # Get each classes' statistics
193         X_i = []
194         for index, result in enumerate(y):
195             if result == label:
196                 X_i.append(X[index])
197         X_i = np.array(X_i)
198         Pi[i] = len(X_i) / N
199         mu[i, :] = np.mean(X_i, axis=0)
200         sigma[i] = np.cov(X_i.T) + self.a*np.eye(d)
201     self.mu = mu
202     self.sigma = sigma
203     self.Pi = Pi
204
205
206 def predict(self, X):
207     C = len(self.Pi)
208     L = np.zeros((C, len(X)))
209     mu = self.mu
210     sigma = self.sigma
211     Pi = self.Pi
212     for i in range(C):
213         #print('start')
214         L[i, :] = multivariate_normal.logpdf(X, mean=mu[i], cov=sigma[i])
215         L[i, :] += np.log(Pi[i])
216     pred = np.argmax(L, axis=0)
217     return pred
218
219 def accuracy(self, X, y):
220     N = len(y)
221     correct = 0
222     pred = self.predict(X)
223     for i in range(N):
224         if pred[i] == y[i]:
225             correct += 1
226     #print('accuracy', correct/N)
227     return correct/N, pred
228
229 training_size = [100, 200, 500, 1000, 2000, 5000, 10000, 30000, 50000]
230
231 qda_accuracy = np.zeros(len(training_size))
232
233
234 for i, N in enumerate(training_size):
235     qda = QDA(1e-8)
236     qda.fit(mnist_train_X[:N, :], mnist_train_y[:N])
237     qda_accuracy[i], _ = qda.accuracy(mnist_val_X, mnist_val_y)
238
239 qda_accuracy
240
241 plt.figure()
242 plt.plot(training_size, (1-qda_accuracy))
243 plt.title('QDA Error Rate')
244 plt.xlabel('Number of Training Samples')
245 plt.ylabel('Error Rate')
246 #plt.savefig("8_3_b.png")
247 plt.show()
248
249

```

```

250 mnist_lda = LDA()
251 mnist_lda.fit(mnist_train_X, mnist_train_y)
252 train_accuracy, train_prev = mnist_lda.accuracy(mnist_train_X, mnist_train_y)
253 val_accuracy, val_prev = mnist_lda.accuracy(mnist_val_X, mnist_val_y)
254 print(train_accuracy, val_accuracy)
255
256 q_mnist_qda = QDA(1e-8)
257 q_mnist_qda.fit(mnist_train_X, mnist_train_y)
258 train_accuracy, train_prev = q_mnist_qda.accuracy(mnist_train_X, mnist_train_y)
259 val_accuracy, val_prev = q_mnist_qda.accuracy(mnist_val_X, mnist_val_y)
260 print(train_accuracy, val_accuracy)
261
262
263 # Training for LDA.
264 training_size = np.array([100, 200, 500, 1000, 2000, 5000, 10000, 30000, 50000])
265 val_acc_list = np.zeros(len(training_size))
266 val_pred_list = []
267
268 for i, N in enumerate(training_size):
269     clf = LDA()
270     clf.fit(mnist_train_X[:N], mnist_train_y[:N])
271     val_acc_list[i], val_pred = clf.accuracy(mnist_val_X, mnist_val_y)
272     val_pred_list.append(val_pred)
273
274 print(val_acc_list)
275 print(len(val_pred_list), len(val_pred_list[0]))
276
277 # Training for QDA
278 training_size = np.array([100, 200, 500, 1000, 2000, 5000, 10000, 30000, 50000])
279 q_val_acc_list = np.zeros(len(training_size))
280 q_val_pred_list = []
281
282 for i, N in enumerate(training_size):
283     clf = QDA(1e-8)
284     clf.fit(mnist_train_X[:N], mnist_train_y[:N])
285     q_val_acc_list[i], val_pred = clf.accuracy(mnist_val_X, mnist_val_y)
286     q_val_pred_list.append(val_pred)
287
288 print(q_val_acc_list)
289 print(len(q_val_pred_list), len(q_val_pred_list[0]))
290
291 def error_evaluate(prediction, labels, N=10):
292     digit_errors = np.zeros((N))
293     prediction = prediction.reshape(len(prediction), 1)
294     for i in range(N):
295         same = (labels == i)
296         total_digit = sum(same)
297         #print(total_digit)
298         correct_digit = sum(prediction[same] == labels[same])
299         error_rate = 1 - (correct_digit/total_digit)
300         digit_errors[i] = error_rate
301     return digit_errors
302
303
304 lda_0 = error_evaluate(val_pred_list[0], np.array(mnist_val_y)).reshape(10,1)
305 lda_1 = error_evaluate(val_pred_list[1], np.array(mnist_val_y)).reshape(10,1)
306 lda_2 = error_evaluate(val_pred_list[2], np.array(mnist_val_y)).reshape(10,1)
307 lda_3 = error_evaluate(val_pred_list[3], np.array(mnist_val_y)).reshape(10,1)
308 lda_4 = error_evaluate(val_pred_list[4], np.array(mnist_val_y)).reshape(10,1)
309 lda_5 = error_evaluate(val_pred_list[5], np.array(mnist_val_y)).reshape(10,1)
310 lda_6 = error_evaluate(val_pred_list[6], np.array(mnist_val_y)).reshape(10,1)
311 lda_7 = error_evaluate(val_pred_list[7], np.array(mnist_val_y)).reshape(10,1)
312 lda_8 = error_evaluate(val_pred_list[8], np.array(mnist_val_y)).reshape(10,1)
313
314 qda_0 = error_evaluate(q_val_pred_list[0], np.array(mnist_val_y)).reshape(10,1)
315 qda_1 = error_evaluate(q_val_pred_list[1], np.array(mnist_val_y)).reshape(10,1)
316 qda_2 = error_evaluate(q_val_pred_list[2], np.array(mnist_val_y)).reshape(10,1)
317 qda_3 = error_evaluate(q_val_pred_list[3], np.array(mnist_val_y)).reshape(10,1)

```

```

318 qda_4 = error_evaluate(q_val_pred_list[4], np.array(mnist_val_y)).reshape(10,1)
319 qda_5 = error_evaluate(q_val_pred_list[5], np.array(mnist_val_y)).reshape(10,1)
320 qda_6 = error_evaluate(q_val_pred_list[6], np.array(mnist_val_y)).reshape(10,1)
321 qda_7 = error_evaluate(q_val_pred_list[7], np.array(mnist_val_y)).reshape(10,1)
322 qda_8 = error_evaluate(q_val_pred_list[8], np.array(mnist_val_y)).reshape(10,1)
323
324 to_plot = np.concatenate((lda_0,lda_1,lda_2,lda_3,lda_4,
325                          lda_5,lda_6,lda_7,lda_8), axis=1)
326 plt.figure(figsize=(10,8))
327 for i in range(10):
328     plt.plot(training_size, to_plot[i], label='digit ' + str(i))
329     plt.legend()
330     plt.title('LDA Classification, digitwise')
331     plt.xlabel('Number of Training Examples')
332     plt.ylabel('Error Rate')
333
334 #plt.savefig("8_3_d_lda.png")
335
336 to_plot = np.concatenate((qda_0,qda_1,qda_2,qda_3,qda_4,
337                          qda_5,qda_6,qda_7,qda_8), axis=1)
338
339 plt.figure(figsize=(10,8))
340 for i in range(10):
341
342     plt.plot(training_size, to_plot[i], label='Digit ' + str(i))
343     plt.legend()
344     plt.title('QDA Classification, digitwise')
345     plt.xlabel('Number of Training Examples')
346     plt.ylabel('Error Rate')
347
348 plt.savefig("8_3_d_qda.png")
349
350 #### 8.4
351
352 #Calculate the HOG feature
353
354 from skimage.feature import hog
355 def hog_dataset(dataset):
356     list_hog_fd = []
357     for sample in dataset:
358         #print(sample.shape)
359         fd = hog(sample.reshape((28, 28)), orientations=8,
360                 pixels_per_cell=(4, 4), cells_per_block=(1, 1),
361                 visualize=False)
362         list_hog_fd.append(fd)
363     return np.array(list_hog_fd, 'float64')
364
365 hog_train_X = hog_dataset(cn_mnist_X)
366 #hog_val_X = hog_dataset(mnist_val_X)
367 hog_test_X = hog_dataset(cn_mnist_test_X)
368
369 std_train_X= np.array([np.std(cn_mnist_X, axis=1)]).T
370 std_test_X = np.array([np.std(cn_mnist_test_X, axis=1)]).T
371
372
373 final_training = np.concatenate((cn_mnist_X, hog_train_X, std_train_X), axis=1)
374
375 final_test = np.concatenate((cn_mnist_test_X, hog_test_X, std_test_X), axis=1)
376
377
378 from sklearn.model_selection import GridSearchCV
379 from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis as QDA
380
381 qda = QDA()
382 param_grid_1 = {"reg_param" : [0.1, 0.5, 1, 2, 5, 10, 15]}
383
384 gs = GridSearchCV(estimator=qda, param_grid=param_grid_1, scoring='accuracy', cv=3, n_jobs=-1)
385

```

```

386 gs = gs.fit(final_training, mnist_y)
387 print(gs.best_score_)
388 print(gs.best_params_)
389
390
391 param_grid_2 = {"reg_param" : [0.01, 0.05, 0.1]}
392
393 gs = GridSearchCV(estimator=qda, param_grid=param_grid_2, scoring='accuracy', cv=3, n_jobs=-1)
394
395 gs = gs.fit(final_training, mnist_y)
396 print(gs.best_score_)
397 print(gs.best_params_)
398
399 param_grid_3 = {"reg_param" : [0.001, 0.01]}
400
401 gs = GridSearchCV(estimator=qda, param_grid=param_grid_3, scoring='accuracy', cv=3, n_jobs=-1)
402
403 gs = gs.fit(final_training, mnist_y)
404 print(gs.best_score_)
405 print(gs.best_params_)
406
407
408 ### 8.5
409
410 spam = io.loadmat("data/spam_data.mat")
411 spam_X, spam_y = spam['training_data'].astype(float), spam['training_labels']
412 spam_test_x = spam['test_data'].astype(float)
413
414
415 spam_train_x, spam_val_x, spam_train_y, spam_val_y = split_train_val_data(spam_X, spam_y, 500)
416
417
418
419 # Computes prior, mean, and covariance
420 def spam_train_lda(train_x, train_y):
421     n, d = train_x.shape
422     priors = np.zeros((2, 1))
423     covariance = np.zeros((d, d)).astype(np.float32)
424     mean_samples = np.zeros((2, d))
425     train_y = train_y.reshape(-1,)
426
427     for spam in [0, 1]:
428         ## Compute priors
429         prob = (train_y == spam).astype(np.int32).sum() / n
430         priors[int(spam)] = prob
431
432         data = train_x[train_y == spam, :]
433         mean_samples[int(spam), :] = data.mean(axis=0)
434         covariance += np.cov(data.T)
435
436     covariance /= 2
437     return priors, mean_samples, covariance
438
439 # Train
440 lda_priors, lda_means, lda_covariance = spam_train_lda(spam_train_x, spam_train_y)
441
442 N_train = spam_train_x.shape[0]
443 out_train = np.zeros((N_train, 2))
444
445 N_val = spam_val_x.shape[0]
446 out_val = np.zeros((N_val, 2))
447
448 for class_y in [0, 1]:
449     prior = lda_priors[class_y]
450     mean = lda_means[class_y, :].reshape(-1, 1)
451
452     w = np.linalg.solve(lda_covariance.T, mean)
453     alpha = -0.5 * w.T.dot(mean) + np.log(prior)

```



```

454
455     out_train[:, class_y] = (spam_train_x.dot(w) + alpha).reshape(-1,)
456     out_val[:, class_y] = (spam_val_x.dot(w) + alpha).reshape(-1,)
457
458 train_pred = np.argmax(out_train, axis=1).reshape(-1, 1)
459 val_pred = np.argmax(out_val, axis=1).reshape(-1, 1)
460
461 def computeDiagonals(X, X_T):
462     N,D = X.shape
463     out_diag = np.zeros((N,1))
464
465     for n in range(N):
466         i = X[n, :].reshape(1, -1)
467         i_t = X_T[:, n].reshape(-1, 1)
468         out_diag[n] = i.dot(i_t)
469
470     return out_diag
471
472 def spam_train_qda(train_x, train_y):
473     n,d = train_x.shape
474     priors = np.zeros((2,1))
475     covariances = np.zeros((2,d,d)).astype(np.float32)
476     mean_samples = np.zeros((2, d))
477     train_y = train_y.reshape(-1,)
478
479     for spam in [0,1]:
480         ## Compute priors
481         prob = (train_y == spam).astype(np.int32).sum() / n
482         priors[int(spam)] = prob
483
484         data = train_x[train_y == spam, :]
485
486         mean_samples[int(spam),:] = data.mean(axis=0)
487         covariances[int(spam), :, :] = np.cov(data.T)
488
489     return priors, mean_samples, covariances
490
491 # Train
492 qda_priors, qda_means, qda_covariances = spam_train_qda(spam_train_x, spam_train_y)
493
494 N_train = spam_train_x.shape[0]
495 out_train = np.zeros((N_train, 2))
496
497 N_val = spam_val_x.shape[0]
498 out_val = np.zeros((N_val, 2))
499
500 for spam in [0,1]:
501     prior = qda_priors[spam]
502     mean = qda_means[spam, :]
503     covariance = qda_covariances[spam, :, :]
504
505     u, sig, v = np.linalg.svd(covariance)
506     covariance = covariance + (sig.min()*(10**-1)) * np.eye(*covariance.shape)
507
508     alpha = -0.5 * np.linalg.det(covariance) + prior
509
510     mean_centered_train = spam_train_x - mean
511     weight_train = np.linalg.solve(covariance, mean_centered_train.T)
512     prediction_train = -0.5*computeDiagonals(mean_centered_train, weight_train) + alpha
513     out_train[:, spam] = prediction_train.reshape(N_train,)
514
515     mean_centered_val = spam_val_x - mean
516     weight_val = np.linalg.solve(covariance, mean_centered_val.T)
517     prediction_val = -0.5*computeDiagonals(mean_centered_val, weight_val) + alpha
518     out_val[:, spam] = prediction_val.reshape(N_val,)
519
520
521 train_pred = np.argmax(out_train, axis=1).reshape(-1, 1)

```

```

522 val_pred = np.argmax(out_val, axis=1).reshape(-1, 1)
523
524
525 from save_csv import results_to_csv
526 N_test = spam_test_x.shape[0]
527 out_test = np.zeros((N_test, 2))
528
529 for class_y in [0,1]:
530     prior = lda_priors[class_y]
531     mean = lda_means[class_y, :].reshape(-1, 1)
532
533     w = np.linalg.solve(lda_covariance.T, mean)
534     alpha = -0.5 * w.T.dot(mean) + np.log(prior)
535
536     out_test[:, class_y] = (spam_test_x.dot(w) + alpha).reshape(-1,)
537
538 test_pred = np.argmax(out_test, axis=1).reshape(-1,)
539 results_to_csv(test_pred)

```